

## Machin Learning Assignment No1s- dataset

Name = nisha sunil ambike, Class :BE , Div:A , Roll No:2

Title: Predict the price of the Uber ride from a given pickup point to the agreed drop-off location. Perform following tasks: 1. Pre-process the dataset. 2. Identify outliers. 3. Check the correlation. 4. Implement linear regression and random forest regression models. Evaluate the models and compare their respective scores like R2, RMSE, etc. Dataset link: <https://www.kaggle.com/datasets/yasserh/uber-fares-dataset>


```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```
In [2]: df = pd.read_csv("C:/Users/Pratibha/Downloads/archive.zip")
df
```

```
Out[2]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	4
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	4
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	4
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	4
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	4
...	...	...	...	...	...	...	...	...
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	40.739367	-73.986525	4
199996	16382965	2014-03-14 01:09:00.0000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	40.736837	-74.006672	4
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	40.756487	-73.858957	4
199998	20259894	2015-05-20 14:56:25.0000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	40.725452	-73.983215	4
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	40.720077	-73.985508	4

200000 rows × 9 columns



1. Pre-process the dataset.

```
In [3]: df.shape
```

```
Out[3]: (200000, 9)
```

```
In [4]: df.head
```

```
Out[4]: <bound method NDFrame.head of Unnamed: 0 key fare_amount \
0 24238194 2015-05-07 19:52:06.0000003 7.5
1 27835199 2009-07-17 20:04:56.0000002 7.7
2 44984355 2009-08-24 21:45:00.00000061 12.9
3 25894730 2009-06-26 08:22:21.0000001 5.3
4 17610152 2014-08-28 17:47:00.000000188 16.0
...
199995 42598914 2012-10-28 10:49:00.00000053 3.0
199996 16382965 2014-03-14 01:09:00.0000008 7.5
199997 27804658 2009-06-29 00:42:00.00000078 30.9
199998 20259894 2015-05-20 14:56:25.0000004 14.5
199999 11951496 2010-05-15 04:08:00.00000076 14.1

pickup_datetime pickup_longitude pickup_latitude \
0 2015-05-07 19:52:06 UTC -73.999817 40.738354
1 2009-07-17 20:04:56 UTC -73.994355 40.728225
2 2009-08-24 21:45:00 UTC -74.005043 40.740770
3 2009-06-26 08:22:21 UTC -73.976124 40.790844
4 2014-08-28 17:47:00 UTC -73.925023 40.744085
...
199995 2012-10-28 10:49:00 UTC -73.987042 40.739367
199996 2014-03-14 01:09:00 UTC -73.984722 40.736837
199997 2009-06-29 00:42:00 UTC -73.986017 40.756487
199998 2015-05-20 14:56:25 UTC -73.997124 40.725452
199999 2010-05-15 04:08:00 UTC -73.984395 40.720077

dropoff_longitude dropoff_latitude passenger_count
0 -73.999512 40.723217 1
1 -73.994710 40.750325 1
2 -73.962565 40.772647 1
3 -73.965316 40.803349 3
4 -73.973082 40.761247 5
...
199995 -73.986525 40.740297 1
199996 -74.006672 40.739620 1
199997 -73.858957 40.692588 2
199998 -73.983215 40.695415 1
199999 -73.985508 40.768793 1

[200000 rows x 9 columns]>
```

```
In [5]: df.isnull()
```

```
Out[5]: Unnamed: 0 key fare_amount pickup_datetime pickup_longitude pickup_latitude dropoff_longitude dropoff_latitude pas
0 False False False False False False False False
1 False False False False False False False False
2 False False False False False False False False
3 False False False False False False False False
4 False False False False False False False False
...
199995 False False False False False False False False
199996 False False False False False False False False
199997 False False False False False False False False
199998 False False False False False False False False
199999 False False False False False False False False

200000 rows x 9 columns
```

```
In [6]: df.drop(columns=["Unnamed: 0", "key"], inplace=True)
df.head()
```

```
Out[6]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.738354	-73.999512	40.723217	1
1	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.728225	-73.994710	40.750325	1
2	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.740770	-73.962565	40.772647	1
3	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	-73.965316	40.803349	3
4	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.744085	-73.973082	40.761247	5

```
In [7]: df.isnull().sum()
```

```
Out[7]: fare_amount      0
pickup_datetime      0
pickup_longitude     0
pickup_latitude      0
dropoff_longitude     1
dropoff_latitude      1
passenger_count      0
dtype: int64
```

```
In [8]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(),inplace = True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].median(),inplace = True)
df.dtypes
```

```
Out[8]: fare_amount      float64
pickup_datetime      object
pickup_longitude     float64
pickup_latitude      float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
In [9]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime)
df.dtypes
```

```
Out[9]: fare_amount      float64
pickup_datetime      datetime64[ns, UTC]
pickup_longitude     float64
pickup_latitude      float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count      int64
dtype: object
```

```
In [10]: df = df.assign(hour = df.pickup_datetime.dt.hour,
day = df.pickup_datetime.dt.day,
month = df.pickup_datetime.dt.month,
year = df.pickup_datetime.dt.year,
dayofweek = df.pickup_datetime.dt.dayofweek)
df
```

Out[10]:

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.999512	40.723217	1	1
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.994710	40.750325	1	2
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.962565	40.772647	1	2
3	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.790844	-73.965316	40.803349	3	
4	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.744085	-73.973082	40.761247	5	1
...	...	...	...	...	...	...	...	...
199995	3.0	2012-10-28 10:49:00+00:00	-73.987042	40.739367	-73.986525	40.740297	1	1
199996	7.5	2014-03-14 01:09:00+00:00	-73.984722	40.736837	-74.006672	40.739620	1	
199997	30.9	2009-06-29 00:42:00+00:00	-73.986017	40.756487	-73.858957	40.692588	2	
199998	14.5	2015-05-20 14:56:25+00:00	-73.997124	40.725452	-73.983215	40.695415	1	1
199999	14.1	2010-05-15 04:08:00+00:00	-73.984395	40.720077	-73.985508	40.768793	1	

200000 rows × 12 columns

In [11]:

df = df.drop(["pickup\_datetime"], axis =1)  
df

Out[11]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	year
0	7.5	-73.999817	40.738354	-73.999512	40.723217	1	19	7	5	2
1	7.7	-73.994355	40.728225	-73.994710	40.750325	1	20	17	7	2
2	12.9	-74.005043	40.740770	-73.962565	40.772647	1	21	24	8	2
3	5.3	-73.976124	40.790844	-73.965316	40.803349	3	8	26	6	2
4	16.0	-73.925023	40.744085	-73.973082	40.761247	5	17	28	8	2
...	...	...	...	...	...	...	...	...	...	...
199995	3.0	-73.987042	40.739367	-73.986525	40.740297	1	10	28	10	2
199996	7.5	-73.984722	40.736837	-74.006672	40.739620	1	1	14	3	2
199997	30.9	-73.986017	40.756487	-73.858957	40.692588	2	0	29	6	2
199998	14.5	-73.997124	40.725452	-73.983215	40.695415	1	14	20	5	2
199999	14.1	-73.984395	40.720077	-73.985508	40.768793	1	4	15	5	2

200000 rows × 11 columns

In [15]:


```
from math import *  
  
def distance_formula(longitude1, latitude1, longitude2, latitude2):  
    travel_dist = []  
  
    for pos in range (len(longitude1)):  
        lon1, lan1, lon2, lan2 = map(radians, [longitude1[pos], latitude1[pos], longitude2[pos], latitude2[pos]  
        dist_lon = lon2 - lon1  
        dist_lan = lan2 - lan1  
  
        a = sin(dist_lan/2)**2 + cos(lan1) * cos(lan2) * sin(dist_lon/2)**2  
  
        #radius of earth = 6371  
        c = 2 * asin(sqrt(a)) * 6371  
        travel_dist.append(c)  
  
    return travel_dist
```

In [16]:

```
df['dist_travel_km'] = distance_formula(df.pickup_longitude.to_numpy(), df.pickup_latitude.to_numpy(), df.dropo
```

2. Identify outliers.

```
In [17]: df.plot(kind = "box",subplots = True,layout = (6,2),figsize=(15,20)) #Boxplot to check the outliers
plt.show()
```



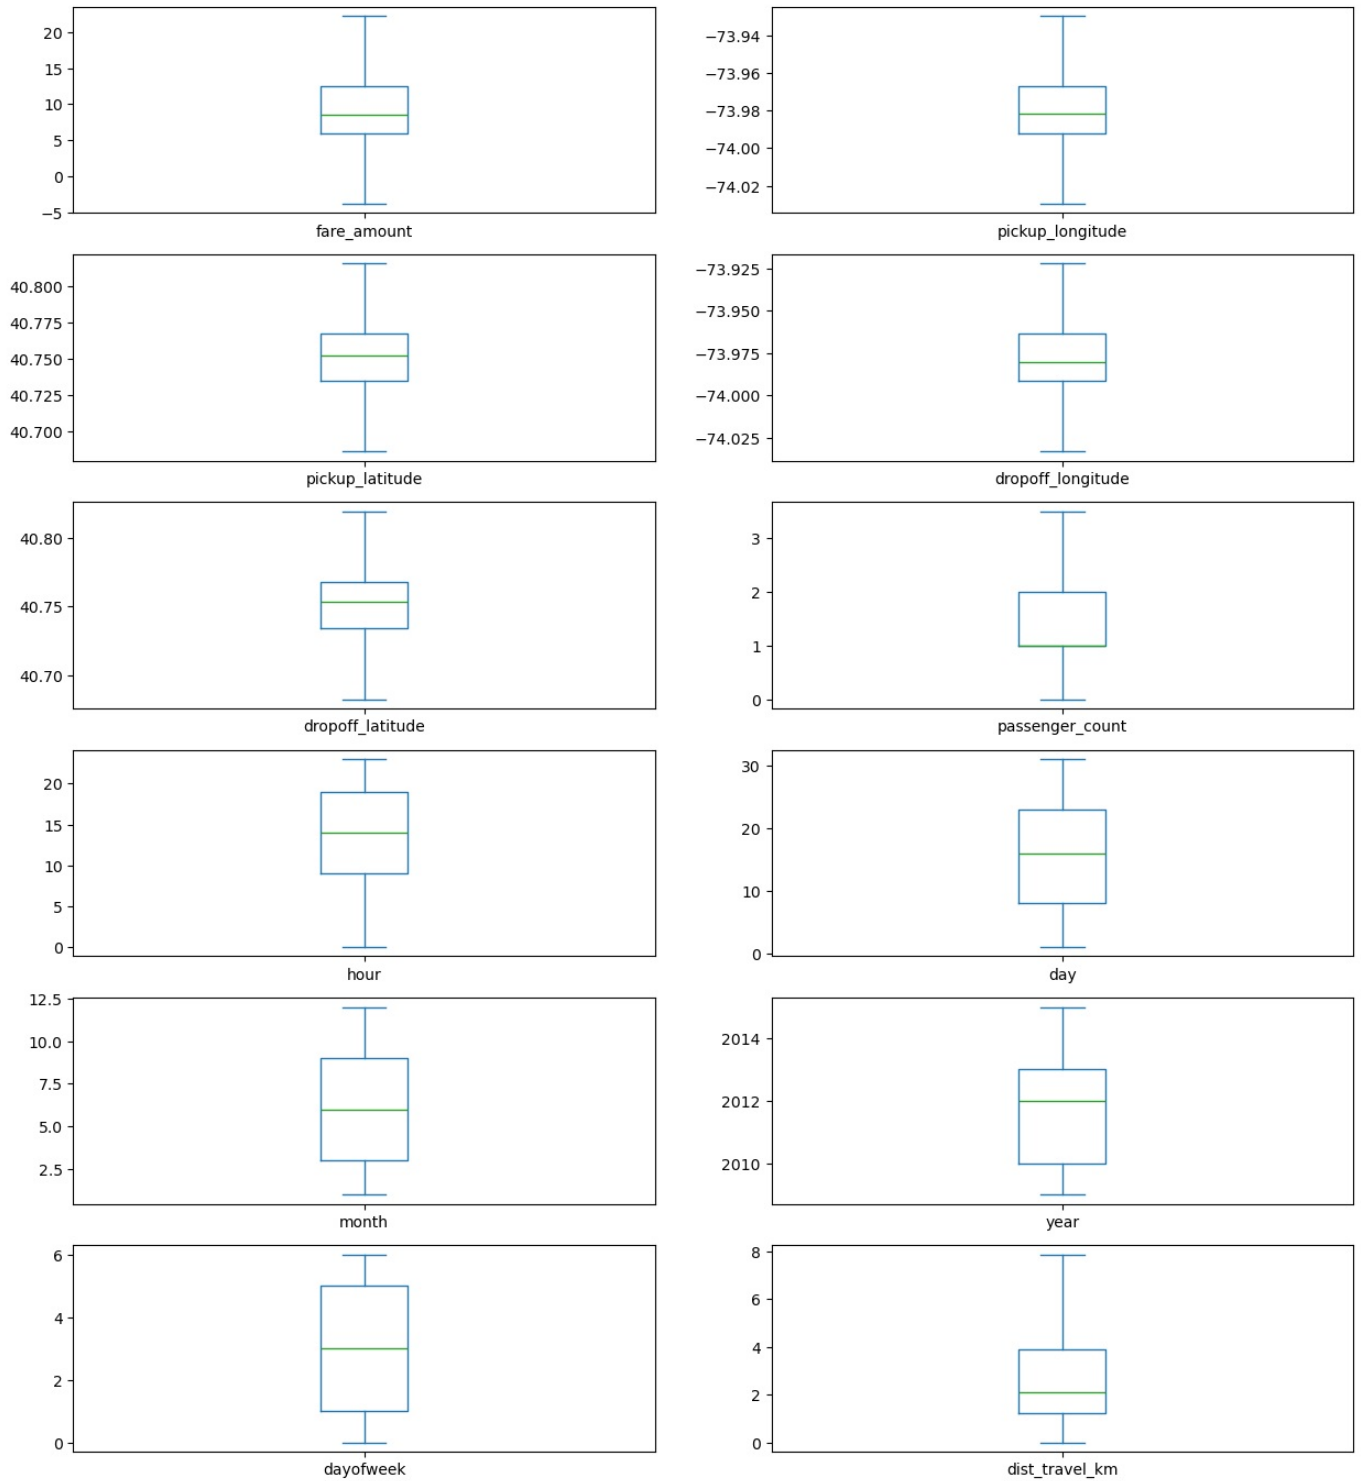
```
In [18]: def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1-1.5*IQR
    upper_whisker = Q3+1.5*IQR
    df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1

def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
```

```
return df1
```

```
In [19]: df = treat_outliers_all(df , df.iloc[:, 0::])
```

```
In [20]: df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15,20))  
plt.show()
```



3. Check the correlation.


```
In [21]: corr = df.corr()  
corr
```

Out[21]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour
fare_amount	1.000000	0.154069	-0.110842	0.218675	-0.125898	0.015778	-0.023623
pickup_longitude	0.154069	1.000000	0.259497	0.425619	0.073290	-0.013213	0.011579
pickup_latitude	-0.110842	0.259497	1.000000	0.048889	0.515714	-0.012889	0.029681
dropoff_longitude	0.218675	0.425619	0.048889	1.000000	0.245667	-0.009303	-0.046558
dropoff_latitude	-0.125898	0.073290	0.515714	0.245667	1.000000	-0.006308	0.019783
passenger_count	0.015778	-0.013213	-0.012889	-0.009303	-0.006308	1.000000	0.020274
hour	-0.023623	0.011579	0.029681	-0.046558	0.019783	0.020274	1.000000
day	0.004534	-0.003204	-0.001553	-0.004007	-0.003479	0.002712	0.004677
month	0.030817	0.001169	0.001562	0.002391	-0.001193	0.010351	-0.003926
year	0.141277	0.010198	-0.014243	0.011346	-0.009603	-0.009749	0.002156
dayofweek	0.013652	-0.024652	-0.042310	-0.003336	-0.031919	0.048550	-0.086947
dist_travel_km	0.844374	0.098094	-0.046812	0.186531	-0.038900	0.009709	-0.038366

```
In [22]: fig,axis = plt.subplots(figsize = (10,6))
sns.heatmap(df.corr(),annot = True) #Correlation Heatmap (Light values means highly correlated)
```

Out[22]: <Axes: >



4. Implement linear regression and random forest regression models.

```
In [23]: df_x = df[['pickup_longitude','pickup_latitude','dropoff_longitude','dropoff_latitude','passenger_count','hour']
df_y = df['fare_amount']
```

```
In [24]: x_train, x_test, y_train, y_test = train_test_split(df_x, df_y, test_size=0.2, random_state=1)
df
```

Out[24]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour	day	month	y
0	7.50	-73.999817	40.738354	-73.999512	40.723217	1.0	19	7	5	2
1	7.70	-73.994355	40.728225	-73.994710	40.750325	1.0	20	17	7	2
2	12.90	-74.005043	40.740770	-73.962565	40.772647	1.0	21	24	8	2
3	5.30	-73.976124	40.790844	-73.965316	40.803349	3.0	8	26	6	2
4	16.00	-73.929786	40.744085	-73.973082	40.761247	3.5	17	28	8	2
...	...	...	...	...	...	...	...	...	...	...
199995	3.00	-73.987042	40.739367	-73.986525	40.740297	1.0	10	28	10	2
199996	7.50	-73.984722	40.736837	-74.006672	40.739620	1.0	1	14	3	2
199997	22.25	-73.986017	40.756487	-73.922036	40.692588	2.0	0	29	6	2
199998	14.50	-73.997124	40.725452	-73.983215	40.695415	1.0	14	20	5	2
199999	14.10	-73.984395	40.720077	-73.985508	40.768793	1.0	4	15	5	2

200000 rows × 12 columns



In [25]:

```
from sklearn.linear_model import LinearRegression

# initialize the linear regression model
reg = LinearRegression()

# Train the model with our training data
reg.fit(x_train, y_train)
```

Out[25]:

▼ LinearRegression

LinearRegression()

In [26]:

```
y_pred_lin = reg.predict(x_test)
print(y_pred_lin)
```

[ 6.27615184 5.09986098 9.43641238 ... 11.07663949 12.15392248
11.41496075]

In [27]:

```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(n_estimators=100)
rf.fit(x_train,y_train)
```

Out[27]:

▼ RandomForestRegressor

RandomForestRegressor()

In [28]:

```
y_pred_rf = rf.predict(x_test)
print(y_pred_rf)
```

[ 4.98 6.5285 9.25 ... 11.5275 11.376 13.13 ]